

Are residual-SAE features less “trivial”?

A two-day investigation into whether training a sparse autoencoder on the *residual* (activation minus a linear prediction from the token embedding) produces fewer single-token features than a normal SAE — and a methodological trap that made the answer look different depending on where I looked.

Notebook write-up 8 tests + hybrid follow-up Verdict: residual null → hybrid works

Bottom line up front

1. **Single-token (“trivial”) features are rare — about 1%, not the ~16% the first metric suggested.** The 16% was a measurement artifact of only looking at each feature’s *strongest* activations.
2. **Residualization does not reduce them.** Full and residual SAEs are equal on every metric that isn’t peak-biased.
3. **Every apparent effect I found along the way** (“residual is more trivial,” a ~22% feature “churn”) turned out to be a *frequency* or *peak-sampling* artifact that vanished under proper controls.
4. The real deliverable is the **method**: how peak-bias and frequency confounds can manufacture effects that aren’t there.
5. **Follow-up (a positive result):** training the map *jointly* with the SAE (the “hybrid”) reconstructs *better* (FVU 0.143 → 0.138) *and* makes the dictionary less trivial (single-token ~halved, holding within frequency bins). Joint training is the ingredient the frozen residual lacked.

The setup & the hypothesis


I trained two sparse autoencoders (SAEs) on the layer-13 residual stream of Gemma-2-2b. Both are BatchTopK SAEs, 16,384 features, k=64, trained on 50M tokens with identical settings. The *only* difference is the input:

- **Full SAE** — trained on the raw activation h .
- **Residual SAE** — trained on $r = h - P(\text{embedding})$, where P is a linear map I’d trained to predict the layer-13 activation from the token’s embedding.


The idea: token identity is the “easy,” linearly-predictable part of the activation. If I subtract it out, maybe the SAE stops wasting features on **trivial single-token detectors** (features that just fire on one word) and spends them on higher-level concepts instead.

The linear map explains a good chunk of the activation — here’s how much variance it captures per layer (I used layer 13):

layer 1



0.85



Test R^2 of the linear map (embedding $\rightarrow$ activation), by layer. Layer 13 (highlighted) sits in a sweet spot: enough token signal to remove, enough context left to matter.

Both SAEs reconstruct the true activation about equally well (the residual SAE's error on `r`, scaled back, matches the full SAE's error on `h` at ~ 0.14 FVU), so any difference between them isn't just one fitting better than the other.

How I measure “trivial”

A feature is **trivial** if it just fires on a single token — it's interpretable, but it's only re-detecting a word the embedding already told the model about. To measure it, for each feature I take its top activating examples and compute two things over the words they fire on:

- **modal fraction** — what fraction are the *same* word (high = single-token).
- **distinct words** — how many different words (low = single-token).

Words are normalized first, so `Happy` / `happy` / `happy` count as one. This whole write-up is really the story of getting this one measurement right.

TEST 1 First pass — basic comparison

Did: Computed the trivial-feature fraction for both SAEs on 100k tokens.

Found: Basically identical — 2.7% vs 2.7%. Looked like a flat null.

PROBLEM, FLAGGED BY MENTOR

100k tokens is far too small. A typical SAE feature fires on roughly 1 in 10,000–1,000,000 tokens, so at 100k I was only measuring the *most frequent* features and missing most of the dictionary. (Neuronpedia's dashboards use ~ 4 M tokens.)

TEST 2 Better metric, much more data

Did: Cleaned up the metric (top-20 examples, string normalization, added the distinct-words measure) and scaled to **4M tokens** using a streaming method, so $\sim 99.8\%$ of features got measured.

Found: Still no clear difference in the distributions. If anything, the residual looked *slightly more* trivial in aggregate — 18% vs 16% pure single-word features — the opposite of the hypothesis.

TEST 3 Controlling for firing frequency

Did: Grouped features into firing-frequency bins and compared the two SAEs *within each bin* — because rare features are naturally more trivial, and the two SAEs might have different mixes of rare vs common features.

Found: The “residual is more trivial” from Test 2 was a **mirage**. It came entirely from the residual SAE having *more rare features* (which are intrinsically more trivial), not from its features being more trivial at any given frequency.

FINDING

At *matched* frequency, there is no real difference — the residual is even slightly *less* trivial in most bins. The naive aggregate comparison was misleading; once controlled, it’s a clean null. **Lesson: frequency is a confound that can flip the sign of the result.**

TEST 4 Actually looking at the features

Did: Built a text “dashboard” (like a homemade Neuronpedia) showing, per feature: how often it fires, an activation histogram, which tokens it predicts, and its top examples with context. Then inspected specific features in both SAEs.

“district” feature

...U.S. **District** Court · Los Angeles County **District** Attorney · **District** of Columbia · d

Fires on “district” across totally different meanings (legal, governmental, geographic). It doesn’t care about meaning, just the word. **Genuinely trivial** — a pure token detector. The metric got this one right.

“health” feature

...single-payer **health** care · lost their **health** insurance · universal **health** care · hea

Fires on “health” but specifically in *healthcare-policy* contexts, and it **predicts the token “insurance.”** So it’s a genuinely **meaningful feature** — yet the metric flags it as trivial because its top examples are all one word.

FINDING

“Single-token” ≠ “useless.” The metric flags token-bound features, but some of them are real, useful concepts. Also: both SAEs have the same district and health detectors — so residualization didn’t delete these concepts.

TEST 5 Matching features across the two SAEs

Did: The systematic version of Test 4. For 100 single-token features in one SAE, find the word each fires on, then find the feature that fires most on that word in the *other* SAE, and check if

it's still single-token or has broadened. Did both directions, plus a **control matching each SAE against itself** to measure how much apparent "change" is just noise.

DIRECTION	BECAME MULTI-WORD	SELF-CONTROL (NOISE FLOOR)	REAL EFFECT
full → residual	30%	8%	~22%
residual → full	18%	6%	~12%

Subtracting the self-control noise floor: ~22% of the full SAE's single-token features broaden in the residual, and ~12% of the residual's are "new."

This looked like a real result: **feature churn** — the count of single-token features stays flat, but *which* concepts have them turns over by 12–22%. I was ready to write this up... and then Test 6 undercut it.

THE CATCH — THIS USED THE "PEAK" METRIC

This whole analysis judged triviality from each feature's *strongest* ~20 activations. Test 6 shows that's exactly the wrong place to look.

TEST 6 Looking across activation strength — the trap

Did: Instead of only the top activations, binned each feature's activations by *strength* (strongest fifth, next fifth, ...) and checked triviality in each bin. A feature might be single-token at its peak but fire on many words when it's activating weakly.

Found: exactly that. Two examples, both labeled "single-token" by the top-20 metric:

feature that is robustly single-token

bin 0 (strongest): **order** ×20 · bin 1: **order** ×20 · bin 2: **order** · bin 3: **order** · bin 4

"order" all the way down. A **real** token detector.

feature that only looks single-token at the peak

bin 0: **news** (but already 84 distinct words) · lower bins: rumors, False, Lied, Myths,

Its top-20 are all "news," so the metric said single-token. But it's really a **misinformation / news-credibility concept feature** — broad and meaningful. The peak hid it.

Another one fired on "subscribing / signing / up / for" across bins — a **"sign up / subscribe" concept feature**, again mislabeled single-token from its peak.


THE TRAP

The top-20 metric is **peak-biased**: it labels a feature by its strongest activations, where almost everything looks token-specific. Several "trivial" features are actually the *higher-level* features I was hoping to find.

TEST 7 Fixing the metric — peak vs range vs weighted

Did: Measured triviality three ways over 4M tokens: on the **peak** (strongest firings), on a **uniform** sample of all firings, and on an **activation-weighted** sample (firings counted in proportion to how strongly the feature fired — the principled middle).

The headline: single-token fraction depends almost entirely on where you look



Fraction of features that are pure single-word, same features, measured three ways. An ~80× swing.

METRIC	SINGLE-WORD (FULL)	SINGLE-WORD (RESID)	EFFECTIVE #WORDS (FULL)
Peak (strongest)	15.7%	17.9%	8.1
Weighted (principled)	0.9%	1.5%	13.2
Uniform	0.3%	0.2%	14.5

"Effective #words" = $\exp(\text{entropy})$ of the word distribution — how many words the feature effectively responds to (1 = pure single-word). Out of 20.


FINDING

Even the activation-*weighted* metric lands near the broad end (~13 effective words, ~1% single-word), **not** halfway. So single-token-ness lives *only* at the razor-thin activation peak; the moment you weight in any of a feature's typical firing, it looks broad. **The ~16% was a peak artifact.**

And on the corrected metric, **full vs residual is null** — both ~1% single-word, ~13 effective words. The ~22% churn from Test 5 was a peak-behavior effect that disappears once you stop looking only at peaks. Within frequency bins it's mixed and tiny.

TEST 8 Frequency × complexity, as one heatmap

Did: Took the metric interpretability dashboards actually use — the number of distinct words in a feature's *max*-20 examples — and plotted it against firing frequency as a 2D heatmap, one panel per SAE. Each *column* is normalized, so it shows $P(\#words \mid frequency)$: among features that fire this often, how many words do they spread over? A third panel shows residual – full. This folds the frequency control (Test 3) and the residual comparison into a single picture, with no arbitrary buckets.



Left / middle: P(#words | freq) for the full and residual SAEs (brighter = more mass). Right: residual – full (blue = residual has more mass there, red = full does). Firing frequency on a log x-axis, half-order-of-magnitude bins.

Found: one dominant structure — a diagonal **ridge from bottom-left to top-right**. Rare features (left) are pinned at 1–2 words; common features (right) spread across ~20. So a **feature’s word-count is largely a function of how often it fires**, not of how abstract it is: a feature isn’t “complex” because it encodes a rich concept — much of its word-count is simply “it fired enough times to have seen many words.”

FINDING

This is the frequency confound as a *picture*: complexity tracks frequency, so any triviality comparison that doesn’t hold frequency fixed is mostly re-measuring the frequency mix. And the **difference panel is near-white everywhere** — conditioned on frequency, full and residual have the same word-count distribution at every frequency. The residual null holds. (The faint blue in the rare corner = residual is a hair *more* single-token among rare features; the lone red cell at the far top-right is the ~2% common tail, too sparse to trust.)

Why the residual ever *looked* more trivial

If the two SAEs are identical at matched frequency, where did the aggregate “residual is more trivial” from Test 2 come from? From the **frequency mix itself**. Here is how many alive features each SAE places in each half-order-of-magnitude frequency bin:

FIRING FREQUENCY	FULL	RESID
20 – 63	176	351
63 – 200	542	955
200 – 632	1,123	1,793
632 – 2,000	2,463	2,206
2,000 – 6,325	3,408	2,796
6,325 – 20,000	5,124	4,897
20,000 – 63,246	3,132	2,996
63,246 – 200,000	322	276
> 200,000	59	79
total alive	16,349	16,349

Alive features per half-order-of-magnitude frequency bin. ~80% of each dictionary lives in the 632–63k range (peak bin: 6.3k–20k); the common tail (>63k) is only ~2%.

The tell is the **rare tail**. Below 632 firings the residual holds **3,099 features (19%) vs the full’s 1,841 (11%)** — nearly double in the two rarest bins. Rare features are

intrinsically single-token (they've simply seen fewer contexts), so a dictionary carrying more of them looks more trivial *on average* — even though, feature-for-feature at matched frequency, nothing has changed. That is the whole “residual is more trivial” effect: a **frequency-composition artifact**, not a complexity difference. (Why the residual's dictionary skews rarer is a separate question these tests don't pin down — plausibly, removing the token-embedding prediction frees the SAE from re-encoding common tokens — but the shift in the *mixture* is real while the change in per-feature triviality is not.)

What I actually learned

1. **Trivial single-token features are rare (~1%), not ~16%.** The high number was a peak-sampling bias — features are token-specific at their strongest activation but broad in typical firing.
2. **Residualization doesn't reduce triviality.** Null on every metric that isn't peak-biased, and null within frequency bins.
3. **Every “effect” I found was an artifact.** The aggregate “residual is more trivial” was a frequency-composition artifact (Test 3); the 22% churn was a peak-sampling artifact (Tests 6–7).
4. **“Single-token” is not “useless.”** Some token-bound features (health→insurance, the misinformation feature) are genuinely meaningful, so the framing of “trivial features are bad” needs rethinking anyway.

THE METHODOLOGICAL TAKEAWAY

Two confounds — **firing frequency** and **peak-sampling** — each independently manufactured an apparent effect that vanished under the right control (matched-frequency bins; activation-weighted measurement). The negative result is only trustworthy *because* those controls are in place.

That said, the null is specifically for a *frozen* map. Training the map **jointly** with the SAE tells a very different story — the follow-up below.

FOLLOW-UP · AFTER THE NULL · MENTOR'S SUGGESTION

HYBRID Training the map jointly — where it finally works

The null above is for a *frozen* map: fit the linear map first, freeze it, then train the SAE on the leftover residual. My mentor suggested the obvious next move — **train the map and the SAE together**, end-to-end, so they coordinate. Call it the **hybrid** SAE:

$\hat{h} = \text{map}(\text{token}) + \text{the SAE's 64 sparse features}$ — the map is a dense, *per-token* baseline the sparse features build on, and both are trained to reconstruct h at once.

The map becomes a smart per-token baseline (versus the SAE’s single *global* bias), it costs zero sparsity budget (k stays 64), and — crucially — it’s optimized *with* the SAE rather than greedily on its own.

Result 1 — it reconstructs better (and joint training is the reason)

SYSTEM	FVU ON H	VS FULL
full SAE	0.1426	—
frozen map + residual SAE	0.1445	+0.0019 (worse)
hybrid (jointly-trained map)	0.1376	-0.0050 (better)

Fraction of variance unexplained on the raw activation, 2M tokens, matched dictionary/k. Lower = better.

FINDING


The hybrid beats the plain SAE by ~3.5% relative FVU at the same sparsity. The *frozen* map ties (even slightly loses) — so it’s specifically the **joint training** that helps: the map settles where it leaves the SAE an *easier* residual, instead of minimizing its own error greedily. Cost: ~5M extra *dense* parameters (the map), so the honest claim is “better reconstruction at matched *sparsity*,” not a free lunch.

Result 2 — the dictionary is also less trivial (the bonus)


This is the part that reaches back to the original question. On the max-20 metric, the hybrid’s single-token fraction is **roughly halved**:

SAE	SINGLE-WORD % (PEAK)	DISTINCT WORDS (PEAK)	EFF #WORDS (WEIGHTED)
full	15.7%	9.20	13.2
residual	17.9%	9.25	12.8
hybrid	9.2%	10.77	14.5

And unlike the residual’s aggregate mirage, this **survives the frequency control**: hybrid features are broader than full’s in *every* firing-frequency bin — even though hybrid has *more* rare features (which should bias it toward looking *more* trivial). So it’s a real per-feature effect, not a composition artifact.



P(#words | frequency) for each SAE, plus **HYBRID – FULL** (right). Unlike the near-white resid–full panel from Test 8, here mass visibly drains out of the single-token corner (red, bottom-left) and shifts up into broader features — a real shift.



Every feature weighted equally (no frequency control). Even in this deliberately naive view, hybrid (green) has a visibly smaller single-token spike and more mass at the broad end.

FINDING

Two-for-one: the hybrid reconstructs *better* and its sparse dictionary is *cleaner*, at matched sparsity. This is the original hypothesis (“residualize → fewer trivial features”) finally working — but via **joint training**, which the frozen residual never delivered.

Looking at the features — single-token → concept

The clinching check: for a given concept word, find the feature that fires most on it in each SAE. **Every time, full uses a single-token detector and hybrid uses a broader coherent concept.**

CONCEPT	FULL — DISTINCT (PEAK)	HYBRID — DISTINCT	WHAT HYBRID’S FEATURE IS
district	1	7	federal district court / legal filings
news	1	1 (10 wtd)	news reporting
health	1	5	single-payer / universal healthcare
insurance	1	5	health insurance / coverage
election	1	16	electoral politics
climate	1	1 (10 wtd)	climate change

“Distinct” = unique tokens in the feature’s max-20 examples (1 = pure single-token). Full is single-token on every concept; hybrid is broader on every one.

Three of these, straight from the feature dashboards (highlighted token = where the feature fires):

“health” — full (single-token) vs hybrid (healthcare policy)

full #7909: ...the **health** risks · **health** disparities · **health** insurance · **health** care ·

hybrid #9204: ...single payer **health** care · Single **Payer** · coverage to **everyone** · **unive**

Full fires on the word “health.” Hybrid fires on the **single-payer / universal-coverage** concept — health, payer, universal, everyone, coverage.

“election” – full (single-token) vs hybrid (electoral politics)

full #7111: ...an **election** year · **election** campaign · **election** ads · **election** season · e

hybrid #14002: ...the **ballot** · **incumbent** Mitch McConnell · City Council **race** · running ·

Full fires on “election.” Hybrid spreads across the whole **electoral field** — ballot, incumbent, race, voters, challenger (16 distinct words at peak vs full’s 1).

“district” – full (single-token) vs hybrid (district court)

full #13449: ...**district** boundaries · **district** attorney · **District** of Columbia · financ

hybrid #3852: ...U.S. **District Court** · **Federal** District Court · **federal** court in Illino

Full fires on “district” across all senses. Hybrid specialized into the **federal-district-court / legal-filing** context — the bare token got decomposed into a context concept.

HONEST CAVEATS

(1) The single-token detection didn’t *vanish* — it moved into the dense map. The *sparse dictionary* is cleaner; the model isn’t representing fewer tokens. (2) It’s **word → topic, not word → abstraction** — hybrid features are broader *topical* concepts (health insurance, climate change), not abstract features like “sycophancy.” (3) Purely-lexical tokens can get *absorbed* rather than broadened: hybrid has no clean “district” feature at all — the map ate the bare token signal, and the best match specialized to the legal sense. (4) The FVU gain is modest and buys ~5M dense params.

Bottom line: the frozen residual was a null, but jointly training the map (hybrid) is a genuine two-for-one — better reconstruction *and* a less-trivial dictionary at matched sparsity. The active ingredient is the *coordination* between map and SAE, not the mere subtraction of a per-token vector.

Open questions & where I’d go next

- **Is “fewer single-token features” even the right goal?** The health/insurance and misinformation features suggest token-bound $\neq$ low-value.
- **Try an earlier layer** (e.g. layer 1, $R^2 \approx 0.85$), where residualization removes far more of the token signal — if any effect exists, it should be strongest there.
- **A non-linear map** instead of a linear one, so more token identity is removed from the residual (token identity clearly survives the linear subtraction — the residual SAE still learns a “district” detector).

Gemma-2-2b · L13 residual · 2 SAEs · 16k feats · k=64 · 50M tok · eval on 4M tokens